

training_overview_1638294786.9697845

December 2, 2021

```
[ ]: import pandas as pd
import os
import matplotlib.pyplot as plt

TIMESTAMP = 1638294786.9697845
DATA_DIRECTORY = '../data'


perf_data = pd.read_csv(os.path.join(DATA_DIRECTORY,
    ↪f'{TIMESTAMP}_performance_stats.csv'))
perf_data['timestamp'] = pd.to_datetime(perf_data['timestamp'])
perf_data.head()
```

```
[ ]: iteration          timestamp  iteration_duration \
0         1 1970-01-01 00:00:01.638297318          2119.456653
1         2 1970-01-01 00:00:01.638298375          1056.551436
2         3 1970-01-01 00:00:01.638299467          1092.254802
3         4 1970-01-01 00:00:01.638300572          1105.114981
4         5 1970-01-01 00:00:01.638302916          2344.427264

      training_duration  examples_read_from_queue  length_train_examples
0         17.130143              10              2000
1         34.403910              40              4000
2         51.698785              18              6000
3         69.033993              19              8000
4         84.150328              18             10000
```


```
[ ]: perf_data.plot.bar(x='iteration', y='examples_read_from_queue')
```

```
[ ]: <AxesSubplot:xlabel='iteration'>
```




```
[ ]: perf_data.plot.bar(x='iteration', y='length_train_examples')
```

```
[ ]: <AxesSubplot:xlabel='iteration'>
```



```
[ ]: perf_data.plot.bar(x='iteration', y='iteration_duration')
```

```
[ ]: <AxesSubplot:xlabel='iteration'>
```



```
[ ]: rndm_data = pd.read_csv(os.path.join(DATA_DIRECTORY,
    ↪f'{TIMESTAMP}_random_player_game_stats.csv'))
rndm_data['timestamp'] = pd.to_datetime(perf_data['timestamp'])
rndm_data.set_index('iteration')
rndm_data.head()
```

```
[ ]: iteration      timestamp  wins  losses  draws  \
0          1 1970-01-01 00:00:01.638297318  2.0    2.0    4.0
1          5 1970-01-01 00:00:01.638298375  3.0    2.0    3.0
2         10 1970-01-01 00:00:01.638299467  2.0    3.0    3.0
3         15 1970-01-01 00:00:01.638300572  0.0    1.0    7.0
4         20 1970-01-01 00:00:01.638302916  0.0    1.0    7.0

nnet_cumul_rewards  random_cumul_rewards
0          0.166667          -0.166667
1          0.166667          -0.166667
2         -0.333333           0.333333
3         -0.166667           0.166667
4         -0.333333           0.333333
```

```
[ ]: hrstc_data = pd.read_csv(os.path.join(DATA_DIRECTORY,
    ↪f'{TIMESTAMP}_heuristic_player_game_stats.csv'))
hrstc_data['timestamp'] = pd.to_datetime(perf_data['timestamp'])
hrstc_data.set_index('iteration')
hrstc_data.head()
```


```
[ ]:      iteration      timestamp  wins  losses  draws  \
0          1 1970-01-01 00:00:01.638297318    0.0    8.0    0.0
1          5 1970-01-01 00:00:01.638298375    0.0    8.0    0.0
2         10 1970-01-01 00:00:01.638299467    0.0    5.0    3.0
3         15 1970-01-01 00:00:01.638300572    0.0    6.0    2.0
4         20 1970-01-01 00:00:01.638302916    0.0    8.0    0.0

      nnet_cumul_rewards  random_cumul_rewards
0          -7.833333          7.833333
1          -7.000000          7.000000
2          -4.833333          4.833333
3          -5.833333          5.833333
4          -7.333333          7.333333
```

```
[ ]: rndm_n_games = int(rndm_data['wins'][0] + rndm_data['losses'][0] +
    ↪rndm_data['draws'][0])
hrstc_n_games = int(rndm_data['wins'][0] + rndm_data['losses'][0] +
    ↪rndm_data['draws'][0])

rndm_fraction_won = rndm_data.apply(lambda row: row['wins'] / rndm_n_games,
    ↪axis=1).to_list()
hrstc_fraction_won = hrstc_data.apply(lambda row: row['wins'] / hrstc_n_games,
    ↪axis=1).to_list()

plt.plot(rndm_data['iteration'], rndm_fraction_won, label=f'Random agent',
    ↪(n={rndm_n_games}))
plt.plot(rndm_data['iteration'], hrstc_fraction_won, label=f'Random agent',
    ↪(n={hrstc_n_games}))
plt.xticks(rndm_data['iteration'])
plt.xlabel('Iteration')
plt.ylabel(f'Fraction won')
plt.legend()
plt.show()
```



```
[ ]: plt.plot(rndm_data['iteration'], rndm_data['nnet_cumul_rewards'],
             ↪label=f'Cumul. rew. (vs. Random agent)')
plt.plot(rndm_data['iteration'], hrstc_data['nnet_cumul_rewards'],
         ↪label=f'Cumul. rew. (vs. Heur. Agent)')
plt.xticks(rndm_data['iteration'])
plt.xlabel('Iteration')
plt.ylabel(f'Reward')
plt.legend()
plt.show()
```

