

training_overview_1638355223.9429553

December 4, 2021

```
[ ]: import pandas as pd
import os
import matplotlib.pyplot as plt

TIMESTAMP = 1638355223.9429553
DATA_DIRECTORY = '../data'


perf_data = pd.read_csv(os.path.join(DATA_DIRECTORY,
    ↪f'{TIMESTAMP}_performance_stats.csv'))
perf_data['timestamp'] = pd.to_datetime(perf_data['timestamp'])
perf_data.head()
```

```
[ ]: iteration          timestamp  iteration_duration \
0      1 1970-01-01 00:00:01.638359832          3766.003021
1      2 1970-01-01 00:00:01.638362029          2197.623581
2      3 1970-01-01 00:00:01.638364174          2144.967788
3      4 1970-01-01 00:00:01.638366413          2238.904411
4      5 1970-01-01 00:00:01.638370947          4533.327736

      training_duration  examples_read_from_queue  length_train_examples
0           3.169783              2              402
1          21.752734             42             2402
2          36.442115             21             4402
3          55.739818             20             6402
4          73.305841             19             8402
```

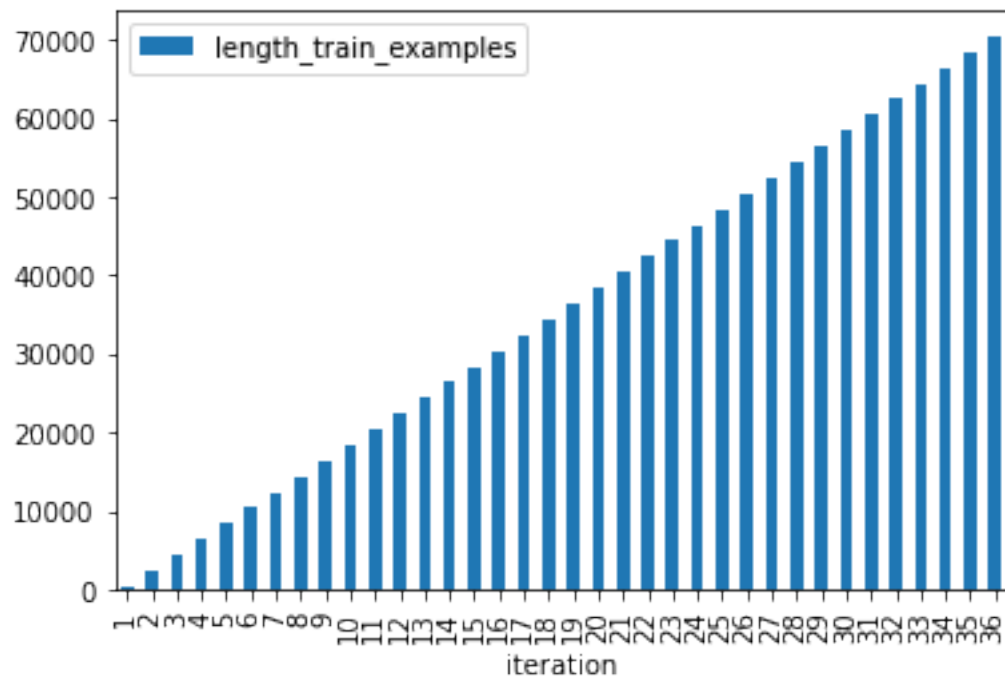
```
[ ]: perf_data.plot.bar(x='iteration', y='examples_read_from_queue')
```

```
[ ]: <AxesSubplot:xlabel='iteration'>
```




```
[ ]: perf_data.plot.bar(x='iteration', y='length_train_examples')
```

```
[ ]: <AxesSubplot:xlabel='iteration'>
```



```
[ ]: perf_data.plot.bar(x='iteration', y='iteration_duration')
```

```
[ ]: <AxesSubplot:xlabel='iteration'>
```



```
[ ]: rndm_data = pd.read_csv(os.path.join(DATA_DIRECTORY,
    ↳ f'{TIMESTAMP}_random_player_game_stats.csv'))
rndm_data['timestamp'] = pd.to_datetime(perf_data['timestamp'])
rndm_data.set_index('iteration')
rndm_data.head()
```

```
[ ]: iteration      timestamp  wins  losses  draws  \
0          1 1970-01-01 00:00:01.638359832    3.0    3.0    2.0
1          5 1970-01-01 00:00:01.638362029    2.0    2.0    4.0
2         10 1970-01-01 00:00:01.638364174    3.0    1.0    4.0
3         15 1970-01-01 00:00:01.638366413    3.0    0.0    5.0
4         20 1970-01-01 00:00:01.638370947    4.0    1.0    3.0
```

```

nnet_cumul_rewards  random_cumul_rewards
0          0.166667          -0.166667
1          0.000000           0.000000
2          0.500000          -0.500000
3          0.666667          -0.666667
4          0.666667          -0.666667
```

```
[ ]: hrstc_data = pd.read_csv(os.path.join(DATA_DIRECTORY,
    ↪f'{TIMESTAMP}_heuristic_player_game_stats.csv'))
hrstc_data['timestamp'] = pd.to_datetime(perf_data['timestamp'])
hrstc_data.set_index('iteration')
hrstc_data.head()
```

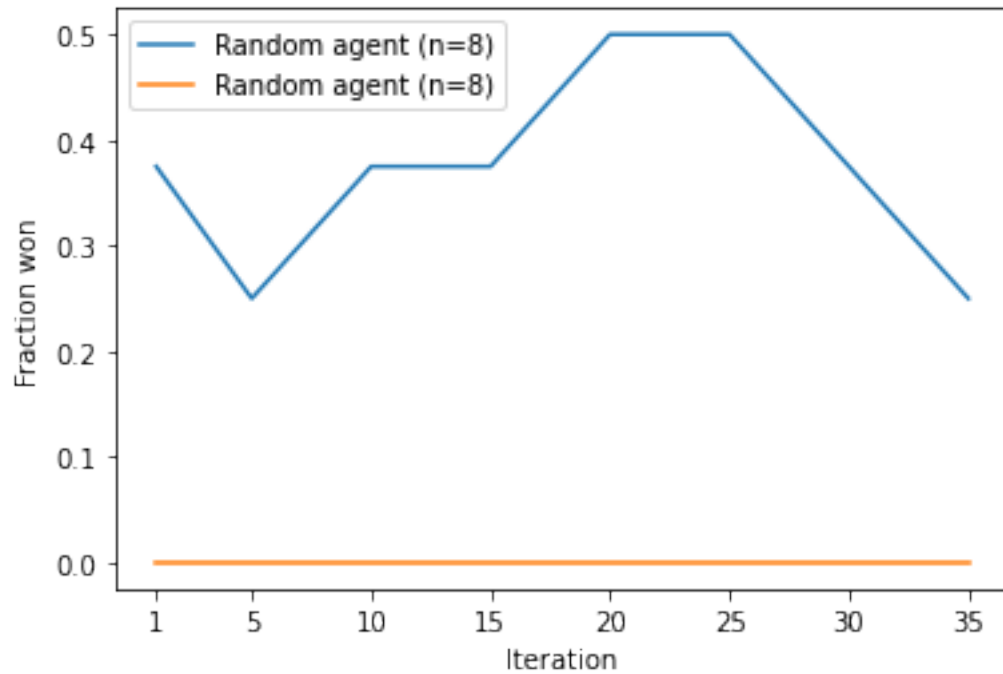
```
[ ]:
iteration      timestamp  wins  losses  draws  \
0           1 1970-01-01 00:00:01.638359832  0.0    8.0    0.0
1           5 1970-01-01 00:00:01.638362029  0.0    7.0    1.0
2          10 1970-01-01 00:00:01.638364174  0.0    8.0    0.0
3          15 1970-01-01 00:00:01.638366413  0.0    8.0    0.0
4          20 1970-01-01 00:00:01.638370947  0.0    8.0    0.0

nnet_cumul_rewards  random_cumul_rewards
0          -8.000000          8.000000
1          -6.833333          6.833333
2          -7.666667          7.666667
3          -8.000000          8.000000
4          -7.333333          7.333333
```

```
[ ]: rndm_n_games = int(rndm_data['wins'][0] + rndm_data['losses'][0] +
    ↪rndm_data['draws'][0])
hrstc_n_games = int(rndm_data['wins'][0] + rndm_data['losses'][0] +
    ↪rndm_data['draws'][0])

rndm_fraction_won = rndm_data.apply(lambda row: row['wins'] / rndm_n_games,
    ↪axis=1).to_list()
hrstc_fraction_won = hrstc_data.apply(lambda row: row['wins'] / hrstc_n_games,
    ↪axis=1).to_list()

plt.plot(rndm_data['iteration'], rndm_fraction_won, label=f'Random agent',
    ↪(n={rndm_n_games}))
plt.plot(rndm_data['iteration'], hrstc_fraction_won, label=f'Random agent',
    ↪(n={hrstc_n_games}))
plt.xticks(rndm_data['iteration'])
plt.xlabel('Iteration')
plt.ylabel(f'Fraction won')
plt.legend()
plt.show()
```



```
[ ]: plt.plot(rndm_data['iteration'], rndm_data['nnet_cumul_rewards'],
             ↪label=f'Cumul. rew. (vs. Random agent)')
plt.plot(rndm_data['iteration'], hrstc_data['nnet_cumul_rewards'],
             ↪label=f'Cumul. rew. (vs. Heur. Agent)')
plt.xticks(rndm_data['iteration'])
plt.xlabel('Iteration')
plt.ylabel(f'Reward')
plt.legend()
plt.show()
```

