

training__overview__1643275233.5268605

February 1, 2022

```
[ ]: import pandas as pd
import os
import matplotlib.pyplot as plt

TIMESTAMP = '1643275233.5268605'
DATA_DIRECTORY = '/run/media/ture/Backup Plus/data/2022-01-28_server_training/'


perf_data = pd.read_csv(os.path.join(DATA_DIRECTORY,
    ↪f'{TIMESTAMP}_performance_stats.csv'))
perf_data['timestamp'] = pd.to_datetime(perf_data['timestamp'])
perf_data.head()
```

```
[ ]: iteration          timestamp  iteration_duration \
0      1 1970-01-01 00:00:01.643277282          1515.136067
1      2 1970-01-01 00:00:01.643278455          1173.027582
2      3 1970-01-01 00:00:01.643280007          1552.738886
3      4 1970-01-01 00:00:01.643281921          1913.236352
4      5 1970-01-01 00:00:01.643285460          3539.642873

      training_duration  examples_read_from_queue  length_experience_buffer
0           89.194729           12           28884
1          470.165958           51          151641
2          781.792642           42          252735
3         1174.927254           53          380306
4         1665.939276           66          539168
```

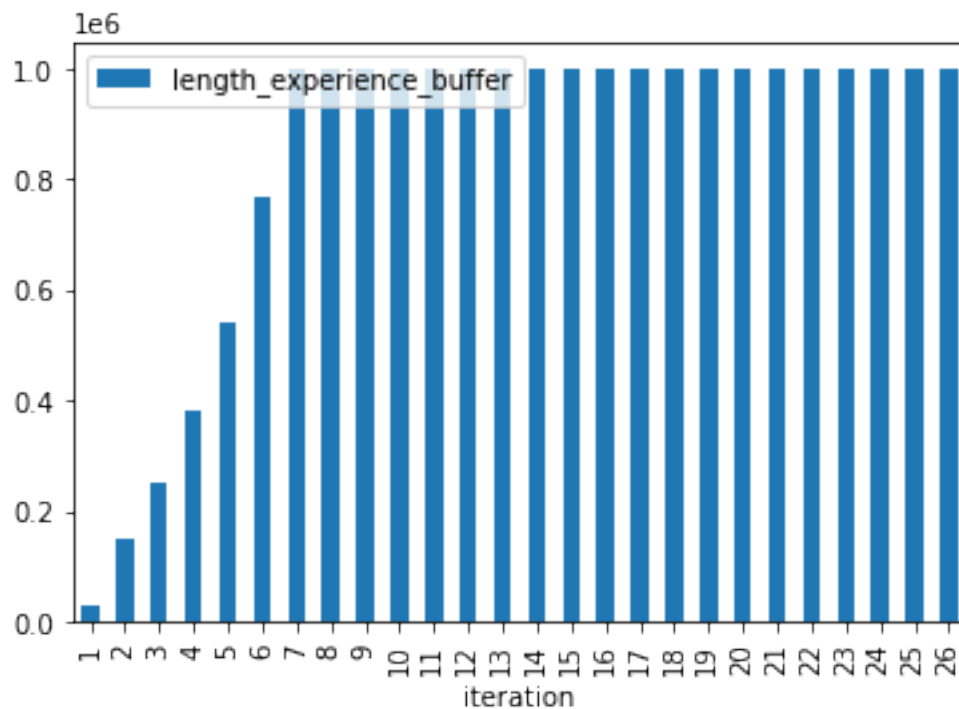
```
[ ]: perf_data.plot.bar(x='iteration', y='examples_read_from_queue')
```

```
[ ]: <AxesSubplot:xlabel='iteration'>
```




```
[ ]: perf_data.plot.bar(x='iteration', y='length_experience_buffer')
```

```
[ ]: <AxesSubplot:xlabel='iteration'>
```



```
[ ]: perf_data.plot.bar(x='iteration', y='iteration_duration')
```

```
[ ]: <AxesSubplot:xlabel='iteration'>
```



```
[ ]: rndm_data = pd.read_csv(os.path.join(DATA_DIRECTORY,
    ↳ f'{TIMESTAMP}_random_player_game_stats.csv'))
rndm_data['timestamp'] = pd.to_datetime(perf_data['timestamp'])
rndm_data.set_index('iteration')
rndm_data.head()
```

```
[ ]: iteration          timestamp  wins  losses  draws  \
0         1 1970-01-01 00:00:01.643277282    4.0    1.0   11.0
1         5 1970-01-01 00:00:01.643278455    4.0    1.0   11.0
2        10 1970-01-01 00:00:01.643280007    3.0    1.0   12.0
3        15 1970-01-01 00:00:01.643281921    3.0    2.0   11.0
4        20 1970-01-01 00:00:01.643285460    1.0    2.0   13.0

nnet_cumul_rewards  random_cumul_rewards
0         1.666667         -1.666667
1         0.666667         -0.666667
2         0.333333         -0.333333
3         0.166667         -0.166667
```

4 -0.166667 0.166667

```
[ ]: hrstc_data = pd.read_csv(os.path.join(DATA_DIRECTORY,
    ↳ f'{TIMESTAMP}_heuristic_player_game_stats.csv'))
hrstc_data['timestamp'] = pd.to_datetime(perf_data['timestamp'])
hrstc_data.set_index('iteration')
hrstc_data.head()
```

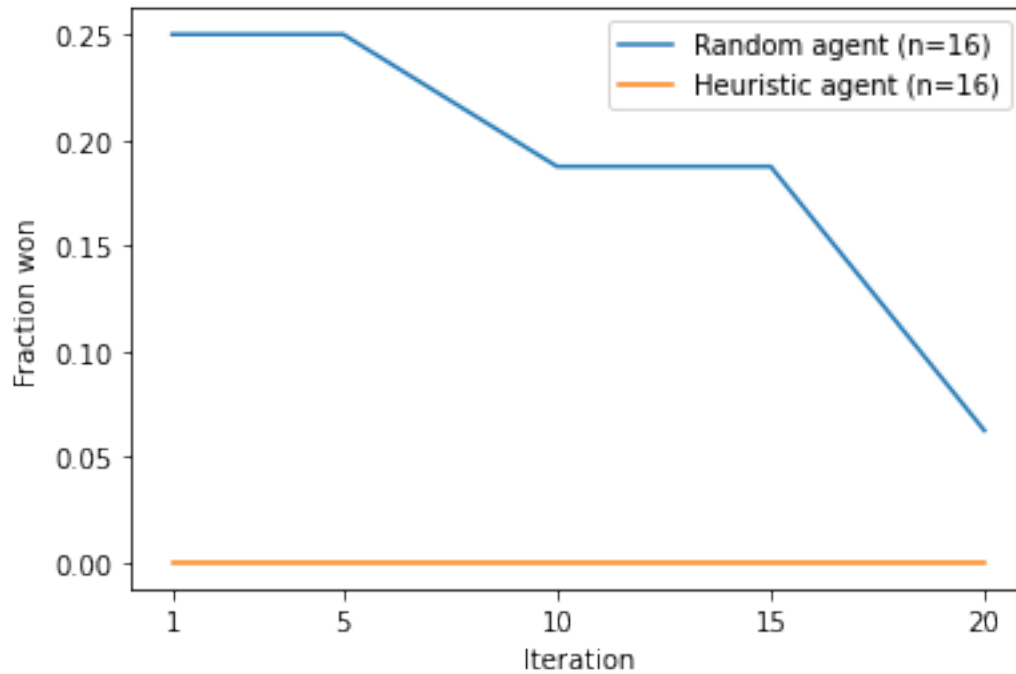
```
[ ]:      iteration      timestamp  wins  losses  draws  \
0         1 1970-01-01 00:00:01.643277282  0.0   16.0    0.0
1         5 1970-01-01 00:00:01.643278455  0.0   10.0    6.0
2        10 1970-01-01 00:00:01.643280007  0.0   11.0    5.0
3        15 1970-01-01 00:00:01.643281921  0.0   13.0    3.0
4        20 1970-01-01 00:00:01.643285460  0.0   14.0    2.0
```

```
      nnet_cumul_rewards  random_cumul_rewards
0         -15.166667      15.166667
1          -8.666667       8.666667
2        -10.833333      10.833333
3        -11.666667      11.666667
4        -12.666667      12.666667
```

```
[ ]: rndm_n_games = int(rndm_data['wins'][0] + rndm_data['losses'][0] +
    ↳ rndm_data['draws'][0])
hrstc_n_games = int(rndm_data['wins'][0] + rndm_data['losses'][0] +
    ↳ rndm_data['draws'][0])

rndm_fraction_won = rndm_data.apply(lambda row: row['wins'] / rndm_n_games,
    ↳ axis=1).to_list()
hrstc_fraction_won = hrstc_data.apply(lambda row: row['wins'] / hrstc_n_games,
    ↳ axis=1).to_list()

plt.plot(rndm_data['iteration'], rndm_fraction_won, label=f'Random agent',
    ↳ (n={rndm_n_games}))
plt.plot(rndm_data['iteration'], hrstc_fraction_won, label=f'Heuristic agent',
    ↳ (n={hrstc_n_games}))
plt.xticks(rndm_data['iteration'])
plt.xlabel('Iteration')
plt.ylabel(f'Fraction won')
plt.legend()
plt.show()
```



```
[ ]: plt.plot(rndm_data['iteration'], rndm_data['nnet_cumul_rewards'],
             ↪label=f'Cumul. rew. (vs. Random agent)')
plt.plot(rndm_data['iteration'], hrstc_data['nnet_cumul_rewards'],
             ↪label=f'Cumul. rew. (vs. Heur. Agent)')
plt.xticks(rndm_data['iteration'])
plt.xlabel('Iteration')
plt.ylabel(f'Reward')
plt.legend()
plt.show()
```

